

### 3. Model building - Random Forest

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import confusion_matrix
# Import RandomForestClassifier model from Scikit-Learn
from sklearn.ensemble import RandomForestClassifier
```

After obtaining cleaning up the data and selecting the variables, we will be trying to build a model to predict whether a stock is worth investing in depending on the column "Class". If class == 1, then it has a positive Price Var and its worth investing in. If class == 0, then it has a negative Price Var and its not worth investing in.

We will be using "Revenue", "EBITDA", "Profit Margin", "returnOnEquity", "Book Value per Share", "Operating Cash Flow", "Dividend Yield", "R&D Expenses" as our predictors.

```
In [2]: df = pd.read_csv('2014_filtered.txt')
df
```

Out[2]:

|      | Revenue      | EBITDA      | Profit<br>Margin | returnOnEquity | Book<br>Value<br>per<br>Share | Operating<br>Cash Flow | Dividend<br>Yield | R&I<br>Expense |
|------|--------------|-------------|------------------|----------------|-------------------------------|------------------------|-------------------|----------------|
| 0    | 5.727000e+09 | 683400000.0 | 0.066            | 0.2041         | 6.706                         | 634100000.0            | 0.0173            | 0.             |
| 1    | 4.551600e+09 | 241900000.0 | -0.021           | -0.1154        | 2.211                         | 536500000.0            | 0.0117            | 0.             |
| 2    | 2.464867e+09 | 771439000.0 | 0.196            | 0.3189         | 3.020                         | 597491000.0            | 0.0000            | 0.             |
| 3    | 3.297600e+09 | 743500000.0 | 0.126            | 0.1969         | 7.779                         | 540300000.0            | 0.0157            | 0.             |
| 4    | 5.973810e+08 | 183876000.0 | 0.122            | 0.1289         | 10.909                        | 111582000.0            | 0.0000            | 0.             |
| ...  | ...          | ...         | ...              | ...            | ...                           | ...                    | ...               | .              |
| 1516 | 1.185080e+08 | 9650000.0   | 0.034            | 0.1324         | 4.249                         | 7612000.0              | 0.0000            | 0.             |
| 1517 | 4.952987e+07 | -53213.0    | -0.002           | -0.0097        | 4.505                         | 523987.0               | 0.0000            | 0.             |
| 1518 | 1.532400e+08 | 20887000.0  | 0.085            | 0.3646         | 2.426                         | -1587000.0             | 0.0000            | 11326000.      |
| 1519 | 3.407580e+08 | 8512000.0   | 0.017            | 0.1456         | 8.489                         | 5745000.0              | 0.0395            | 0.             |
| 1520 | 4.033737e+07 | 4959141.0   | 0.060            | 0.0721         | 1.645                         | 4012331.0              | 0.0000            | 3379920.       |

1521 rows × 11 columns

```
In [3]: y = pd.DataFrame(df["Class"])
```

```

predictors = pd.DataFrame(df[["Revenue", "EBITDA", "Profit Margin", "returnOnEquity",
                             "Operating Cash Flow", "Dividend Yield"]])
X = predictors

```

```

In [6]: rforest = RandomForestClassifier(n_estimators = 500, # n_estimators denote number of
                                         max_depth = 6)      # set the maximum depth of each tree

# Split the Dataset into Train and Test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)

# Random Forest using Train Data
rforest = RandomForestClassifier(n_estimators = 100, max_depth = 4) # create the object
rforest.fit(X_train, y_train) # train the model

# Predict Response corresponding to Predictors
y_train_pred = rforest.predict(X_train)
y_test_pred = rforest.predict(X_test)

#Getting the confusion matrices
train_confusion = confusion_matrix(y_train, y_train_pred)

TP_train = train_confusion[1][1]
FN_train = train_confusion[1][0]
TN_train = train_confusion[0][0]
FP_train = train_confusion[0][1]

test_confusion = confusion_matrix(y_test, y_test_pred)

TP_test = test_confusion[1][1]
FN_test = test_confusion[1][0]
TN_test = test_confusion[0][0]
FP_test = test_confusion[0][1]

# Check the Goodness of Fit (on Train Data)
print("Goodness of Fit of Model \tTrain Dataset")
print("Classification Accuracy \t:", rforest.score(X_train, y_train))
print("True Postitive Rate \t\t:", TP_train/(TP_train+FN_train))
print("True Negative Rate \t\t:", TN_train/(TN_train+FP_train))
print("False Negative Rate \t\t:", FN_train/(FN_train+TP_train))
print("False Postitive Rate \t\t:", FP_train/(FP_train+TN_train))
print()

# Check the Goodness of Fit (on Test Data)
print("Goodness of Fit of Model \tTest Dataset")
print("Classification Accuracy \t:", rforest.score(X_test, y_test))
print("True Postitive Rate \t\t:", TP_test/(TP_test+FN_test))
print("True Negative Rate \t\t:", TN_test/(TN_test+FP_test))
print("False Negative Rate \t\t:", FN_test/(FN_test+TP_test))
print("False Postitive Rate \t\t:", FP_test/(FP_test+TN_test))
print()

# Plot the Confusion Matrix for Train and Test
f, axes = plt.subplots(1, 2, figsize=(12, 4))
sb.heatmap(train_confusion,
           annot = True, fmt=".0f", annot_kws={"size": 18}, ax = axes[0])
sb.heatmap(test_confusion,
           annot = True, fmt=".0f", annot_kws={"size": 18}, ax = axes[1])

```

<ipython-input-6-f933f0314320>:9: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```

rforest.fit(X_train, y_train) # train the model

```


Goodness of Fit of Model  
Classification Accuracy  
True Postitive Rate  
True Negative Rate  
False Negative Rate  
False Postitive Rate

Train Dataset  
: 0.6947368421052632  
: 0.5444646098003629  
: 0.8353140916808149  
: 0.455535390199637  
: 0.16468590831918506

Goodness of Fit of Model  
Classification Accuracy  
True Postitive Rate  
True Negative Rate  
False Negative Rate  
False Postitive Rate

Test Dataset  
: 0.6377952755905512  
: 0.44642857142857145  
: 0.7887323943661971  
: 0.5535714285714286  
: 0.2112676056338028

Out[6]: <AxesSubplot:>



Utilizing the random forest model, we managed to obtain a classification accuracy of around 0.69 for the training data set and a classification accuracy of around 0.63 for the test data set. Both of which are strongly(?) accurate.

The random forest model also has a decent(?) true positive rate of 0.54 for the train and 0.44 for the test data sets. However, the true negative rate is high at 0.83 for the train and 0.78 for the test data set. A better model is needed to improve on the true positive rates. We thus, turn to Grid Search.

In [ ]: